Himabindu Lakkaraju

Contact Information	491 Morgan Hall
	15 Harvard Way
	Boston, MA 02163
	E-mail: hlakkaraju@hbs.edu; hlakkaraju@seas.harvard.edu
	Webpage: http://himalakkaraju.github.io

Research Interests

Trustworthy Machine Learning (Interpretability, Fairness, Robustness, and Privacy); Large Language Models; Human-AI Interaction; Applications of AI/ML to Decision Making in Healthcare, Law, and Policy.

Academic & Professional Experience

Harvard University   *01/2020 - Present*
Assistant Professor with appointments in the Business School and the Department of Computer Science (Affiliate)
Faculty Affiliate, Harvard Data Science Initiative

Simons Institute for the Theory of Computing, UC Berkeley
Visiting Scientist, Summer Cluster on Interpretable Machine Learning   *06/2022 - 08/2022*
Visiting Graduate Student, Summer Cluster on Algorithmic Fairness   *07/2018 - 08/2018*

Fiddler AI   *06/2021 - 11/2022*
Chief AI Research Fellow

Harvard University   *11/2018 - 12/2019*
Postdoctoral Fellow, Business School & Department of Computer Science

Stanford University   *9/2012 - 9/2018*
Research Assistant, Department of Computer Science

Microsoft Research, Redmond
Visiting Researcher   *5/2017 - 6/2017*
Research Intern   *6/2016 - 9/2016*

University of Chicago   *6/2014 - 8/2014*
Data Science for Social Good Fellow

IBM Research   *7/2010 - 7/2012*
Research Engineer

Education

Stanford University   *9/2012 - 9/2018*
Ph.D. in Computer Science

Stanford University   *9/2012 - 9/2015*
Master of Science (MS) in Computer Science

Indian Institute of Science (IISc)   *8/2008 - 7/2010*
Master of Engineering (MEng) in Computer Science & Automation

Selected Honors & Achievements

NSF CAREER Award	2023
AI2050 Early Career Fellowship by Schmidt Futures	2023
Named **Kavli Fellow** by the National Academy of Sciences	2023
Adobe Data Science Research Award	2023
Best Paper Award, **ICML Workshop on Interpretable ML in Healthcare**	2022
Outstanding Paper Award Honorable Mention **NeurIPS Workshop on Trustworthy and Socially Responsible Machine Learning**	2022
JP Morgan Faculty Research Award	2022

Selected as a member of the **National AI Advisory Committee** instituted by the US government (could not serve due to citizenship status)	2022
National Science Foundation (NSF) Amazon Fairness in AI Grant	2021
Google AI for Social Good Research Award	2021
Best Paper Runner Up, **ICML Workshop on Algorithmic Recourse**	2021
Advisory Board Member, **The Stanford Center for Legal Informatics**	2020 - Present
Google Research Award	2020
Amazon Research Award	2020
Co-founded **Trustworthy ML Initiative** with the goal of enabling easy access to resources on trustworthy ML & to build a community of researchers/practitioners	2020
Hoopes Prize for undergraduate thesis mentoring, Harvard University	2020
Named as one of the **35 Innovators Under 35** (Global) by MIT Tech Review	2019
Named as an **Innovator to Watch** by Vanity Fair	2019
Selected for the prestigious **Cowles Fellowship** by Yale University (declined)	2018
INFORMS Data Mining Best Paper Award	2017
Microsoft Research Dissertation Grant	2017
Named as one of the **Rising Stars in Computer Science**	2016
Outstanding Reviewer Award International World Wide Web Conference (WWW)	2016
Google Anita Borg Fellowship in recognition of research and leadership	2015
Stanford Graduate Fellowship for exceptional academic performance Awarded to top 3% of Stanford Ph.D. students	2013-17
Eminence and Excellence Award for outstanding research contributions IBM Research	2012
Best Paper Award, **SIAM International Conference on Data Mining (SDM)**	2011
All India Rank 32 (99.82%ile) Graduate Aptitude Test in Engineering (GATE) Entrance examination for IISc & IITs in Computer Science & Engineering	2008

Selected Grants & Fellowships

As Faculty

NSF CAREER Award (US$550,664) – Sole PI	2023 - 2028
AI2050 Early Career Fellowship by Schmidt Futures (US$300,00) – Sole PI	2023 - 2026
Adobe Data Science Research Award (US$50,000) – Sole PI	2023 - 2024
D3 Institute at Harvard Grant (US$600,000) – Sole PI	2022 - 2025
JP Morgan Faculty Research Award (US$110,000) – Sole PI	2022 - 2024
NSF-Amazon Fairness in AI (FAI) grant (US$375,000) – co-PI	2021 - 2024
Amazon Faculty Research Award (US$70,000) – Sole PI	2021 - 2024
Google AI for Social Good Research Award (US$10,000) – Sole PI	2021 - 2022
Google Research Award (US$600,000) – PI	2020 - 2024
NSF IIS: Robust Intelligence (RI) Small (US$450,000) – Harvard PI	2020 - 2023
Bayer Trust in Science Award (US$100,000) – PI	2020 - 2021

As Student

Microsoft Research Dissertation Grant (US$20,000)	2017

Stanford Graduate Fellowship (tuition + US$41,700 p.a.)	2013 - 2017
Google Anita Borg Scholarship (US$10,000)	2015
Facebook Graduate Fellowship Finalist (US$500)	2013
Indian Institute of Science Graduate Scholarship (tuition + Rs.96,000 p.a.)	2008 - 2010
SAP India Research Grant (Rs.150,000)	2009 - 2010

Research Articles   **Total Citations**: **6634**   **h-index**: **34**   **i10-index**: **46**

(* below indicates equal contribution)

Book Chapters

[71] Analyzing Human Decisions and Machine Predictions in Bail Decision Making
Jon Kleinberg, Himabindu Lakkaraju, Jure Leskovec, Jens Ludwig, Sendhil Mullainathan
(author names are ordered alphabetically)
The Inequality Reader: Contemporary and Foundational Readings in Race, Class, and Gender; Third Edition, 2022.

Articles in Peer-Reviewed Journals

[70] TalkToModel: Explaining Machine Learning Models with Interactive Natural Language Conversations
Dylan Slack, Satyapriya Krishna, Himabindu Lakkaraju*, Sameer Singh*
Nature Machine Intelligence - 2023.
Outstanding Paper Award Honorable Mention, NeurIPS Workshop on Trustworthy and Socially Responsible ML, 2022.

[69] Evaluating Explainability for Graph Neural Networks
Chirag Agarwal, Owen Queen, Himabindu Lakkaraju, Marinka Zitnik
Nature Scientific Data - 2023.

[68] When Does Uncertainty Matter?: Understanding the Impact of Predictive Uncertainty in ML Assisted Decision Making
Sean McGrath, Parth Mehta, Alexandra Zytek, Isaac Lage, Himabindu Lakkaraju
TMLR - Transactions on Machine Learning Research, 2023.
Featured in VentureBeat

[67] Human Decisions and Machine Predictions
Jon Kleinberg, Himabindu Lakkaraju, Jure Leskovec, Jens Ludwig, Sendhil Mullainathan
QJE - Quarterly Journal of Economics, 2018.
(author names are ordered alphabetically)
Featured in MIT Technology Review, Harvard Business Review, The New York Times, and as Research Spotlight on National Bureau of Economics front page

[66] Mining Digital Footprints to Extract Patterns and Predict Real-Life Outcomes
Michal Kosinski, Yilun Wang, Himabindu Lakkaraju, Jure Leskovec
Psychological Methods - 2016.

Articles in Peer-Reviewed Conference Proceedings

[65] Post hoc Explanations of Language Models can Improve Language Models
Satyapriya Krishna, Jiaqi Ma, Dylan Slack, Asma Ghandeharioun, Sameer Singh, Himabindu Lakkaraju
NeurIPS - Advances in Neural Information Processing Systems, 2023.

[64] Verifiable Feature Attributions: A Bridge between Post Hoc Explainability and Inherent Interpretability
Usha Bhalla*, Suraj Srinivas*, Himabindu Lakkaraju
NeurIPS - Advances in Neural Information Processing Systems, 2023.

[63] Which Models have Perceptually-Aligned Gradients? An Explanation via Off-Manifold Robustness
Suraj Srinivas*, Sebastian Bordt*, Himabindu Lakkaraju
NeurIPS - Advances in Neural Information Processing Systems, 2023.
Spotlight Presentation (Top 3%)

[62] M4: A Unified XAI Benchmark for Faithfulness Evaluation of Feature Attribution Methods across Metrics, Modalities, and Models
Xuhong Li, Mengnan Du, Jiamin Chen, Yekun Chai, Himabindu Lakkaraju, Haoyi Xiong
NeurIPS - Advances in Neural Information Processing Systems, 2023.

[61] Towards Bridging the Gaps between the Right to Explanation and the Right to be Forgotten
Satyapriya Krishna*, Jiaqi Ma*, Himabindu Lakkaraju
ICML - International Conference on Machine Learning, 2023.

[60] On the Impact of Actionable Explanations on Social Segregation
Ruijiang Gao, Himabindu Lakkaraju
ICML - International Conference on Machine Learning, 2023.

[59] On Minimizing the Impact of Dataset Shifts on Actionable Explanations
Anna Meyer*, Dan Ley*, Suraj Srinivas, Himabindu Lakkaraju
UAI - Conference on Uncertainty in Artificial Intelligence, 2023.
Oral Presentation (Top 5%)

[58] Probabilistically Robust Recourse: Navigating the Trade-offs between Costs and Robustness in Algorithmic Recourse
Martin Pawelczyk, Teresa Datta, Johannes van den Heuvel, Gjergji Kasneci, Himabindu Lakkaraju
ICLR - International Conference on Learning Representations, 2023.

[57] On the Privacy Risks of Algorithmic Recourse
Martin Pawelczyk, Himabindu Lakkaraju*, Seth Neel*
AISTATS - International Conference on Artificial Intelligence and Statistics, 2023.

[56] Which Explanation Should I Choose? A Function Approximation Perspective to Characterizing Post hoc Explanations
Tessa Han, Suraj Srinivas, Himabindu Lakkaraju
NeurIPS - Advances in Neural Information Processing Systems (NeurIPS), 2022.
Best Paper Award, ICML Workshop on Interpretable ML in Healthcare, 2022.

[55] Flatten the Curve: Efficiently Training Low-Curvature Neural Networks
Suraj Srinivas, Kyle Matoba, Himabindu Lakkaraju, Francois Fleuret
NeurIPS - Advances in Neural Information Processing Systems (NeurIPS), 2022.

[54] OpenXAI: Towards a Transparent Evaluation of Model Explanations
Chirag Agarwal, Satyapriya Krishna, Eshika Saxena, Martin Pawelczyk, Nari Johnson, Isha Puri, Marinka Zitnik, Himabindu Lakkaraju
NeurIPS - Advances in Neural Information Processing Systems (NeurIPS), 2022.

[53] Data Poisoning Attacks on Off-Policy Evaluation Methods
Elita Lobo, Harvineet Singh, Marek Petrik, Cynthia Rudin, Himabindu Lakkaraju
UAI - Conference on Uncertainty in Artificial Intelligence, 2022.
Oral Presentation (Top 5%)

[52] Exploring Counterfactual Explanations Through the Lens of Adversarial Examples: A Theoretical and Empirical Analysis
Martin Pawelczyk, Chirag Agarwal, Shalmali Joshi, Sohini Upadhyay, Himabindu Lakkaraju
AISTATS - International Conference on Artificial Intelligence and Statistics, 2022.

[51] Probing GNN Explainers: A Rigorous Theoretical and Empirical Analysis of GNN Explanation Methods
Chirag Agarwal, Marinka Zitnik*, Himabindu Lakkaraju*
AISTATS - International Conference on Artificial Intelligence and Statistics, 2022.

[50] Fairness via Explanation Quality: Evaluating Disparities in the Quality of Post hoc Explanations
Jessica Dai, Sohini Upadhyay, Ulrich Aivodji, Stephen Bach, Himabindu Lakkaraju
AIES - AAAI/ACM Conference on AI, Ethics, and Society, 2022.

[49] Towards Robust Off-Policy Evaluation via Human Inputs
Harvineet Singh, Shalmali Joshi, Finale Doshi-Velez, Himabindu Lakkaraju
AIES - AAAI/ACM Conference on AI, Ethics, and Society, 2022.

[48] A Human-Centric Perspective on Model Monitoring
Murtuza N Shergadwala, Himabindu Lakkaraju, Krishnaram Kenthapadi
HCOMP - AAAI Conference on Human Computation and Crowdsourcing, 2022.

[47] Towards Robust and Reliable Algorithmic Recourse
Sohini Upadhyay*, Shalmali Joshi*, Himabindu Lakkaraju
NeurIPS - Advances in Neural Information Processing Systems (NeurIPS), 2021.
Best Paper Runner Up, ICML Workshop on Algorithmic Recourse, 2021.

[46] Reliable Post hoc Explanations: Modeling Uncertainty in Explainability
Dylan Slack, Sophie Hilgard, Sameer Singh, Himabindu Lakkaraju
NeurIPS - Advances in Neural Information Processing Systems, 2021.

[45] Counterfactual Explanations Can Be Manipulated
Dylan Slack, Sophie Hilgard, Himabindu Lakkaraju, Sameer Singh
NeurIPS - Advances in Neural Information Processing Systems, 2021.

[44] Learning Models for Algorithmic Recourse
Alexis Ross, Himabindu Lakkaraju, Osbert Bastani
NeurIPS - Advances in Neural Information Processing Systems, 2021.

[43] Towards the Unification and Robustness of Perturbation and Gradient Based Explanations
Sushant Agarwal, Shahin Jabbari, Chirag Agarwal*, Sohini Upadhyay*, Steven Wu, Himabindu Lakkaraju
ICML - International Conference on Machine Learning, 2021.
Shorter version presented at Foundations of Responsible Computing (FORC), 2022.

[42] Towards a Unified Framework for Fair and Stable Graph Representation Learning
Chirag Agarwal, Himabindu Lakkaraju*, Marinka Zitnik*
UAI - Conference on Uncertainty in Artificial Intelligence, 2021.
Oral Presentation (Top 5%)

[41] Does Fair Ranking Improve Minority Outcomes? Understanding the Interplay of Human and Algorithmic Biases in Online Hiring
Tom Suhr, Sophie Hilgard, Himabindu Lakkaraju
AIES - AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society, 2021.

[40] Fair influence maximization: A welfare optimization approach
Aida Rahmattalabi, Shahin Jabbari, Himabindu Lakkaraju, Phebe Vayanos, Eric Rice, Milind Tambe
AAAI - AAAI International Conference on Artificial Intelligence, 2021.

[39] Beyond Individualized Recourse: Interpretable and Interactive Summaries of Actionable Recourses
Kaivalya Rawal, Himabindu Lakkaraju
NeurIPS - Advances in Neural Information Processing Systems, 2020.

[38] Incorporating Interpretable Output Constraints in Bayesian Neural Networks
Wanqian Yang, Lars Lorch, Moritz Gaule, Himabindu Lakkaraju, Finale Doshi-Velez
NeurIPS - Advances in Neural Information Processing Systems, 2020.
Spotlight Presentation (Top 3%)

[37] Robust and Stable Black Box Explanations
Himabindu Lakkaraju, Nino Arsov, Osbert Bastani
ICML - International Conference on Machine Learning, 2020

[36] How do I fool you?: Manipulating User Trust via Misleading Black Box Explanations
Himabindu Lakkaraju, Osbert Bastani
AIES - AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society, 2020.
Oral Presentation (Top 16.6%)

[35] Fooling LIME and SHAP: Adversarial Attacks on Post hoc Explanation Methods
Dylan Slack, Sophie Hilgard, Emily Jia, Sameer Singh, Himabindu Lakkaraju
AIES - AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society, 2020.
Featured in Harvard Business Review and deeplearning.ai
Best Paper (Non-Archival) at AAAI Workshop on Safe AI, 2020
Oral Presentation (Top 16.6%)

[34] Faithful and Customizable Explanations of Black Box Models
Himabindu Lakkaraju, Ece Kamar, Rich Caruana, Jure Leskovec
AIES - AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society, 2019.
Oral Presentation (Top 10%)

[33] The Selective Labels Problem: Evaluating Algorithmic Predictions in the Presence of Unobservables
Himabindu Lakkaraju, Jon Kleinberg, Jure Leskovec, Jens Ludwig, Sendhil Mullainathan
KDD - ACM SIGKDD Conference on Knowledge Discovery and Data Mining, 2017.
Oral Presentation (Top 8.5%)

[32] Learning Cost-Effective and Interpretable Treatment Regimes
Himabindu Lakkaraju, Cynthia Rudin
AISTATS - International Conference on Artificial Intelligence and Statistics, 2017.
INFORMS Data Mining Best Paper Award, 2017

[31] Identifying Unknown-Unknowns in the Open World: Representations and Policies for Guided Exploration
Himabindu Lakkaraju, Ece Kamar, Rich Caruana, Eric Horvitz
AAAI - AAAI International Conference on Artificial Intelligence, 2017.
Featured in Bloomberg Technology

[30] Confusions over Time: An Interpretable Bayesian Model for Characterizing Trends in Decision Making
Himabindu Lakkaraju, Jure Leskovec
NIPS - Advances in Neural Information Processing Systems, 2016.

[29] Interpretable Decision Sets: A Joint Framework for Description and Prediction
Himabindu Lakkaraju, Stephen Bach, Jure Leskovec
KDD - ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, 2016.

[28] A Machine Learning Framework to Identify Students at Risk of Adverse Academic Outcomes
Himabindu Lakkaraju, Everaldo Aguiar, Carl Shan, David Miller, Nasir Bhanpuri, Rayid Ghani, Kecia Addison
KDD - ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, 2015.
Oral Presentation (Top 8.2%)

[27] A Bayesian Framework for Modeling Human Evaluations
Himabindu Lakkaraju, Jure Leskovec, Jon Kleinberg, Sendhil Mullainathan
SDM - SIAM International Conference on Data Mining, 2015.
Oral Presentation (Top 5%)

[26] Who, When, and Why: A Machine Learning Approach to Prioritizing Students at Risk of not Graduating High School on Time
Everaldo Aguiar, Himabindu Lakkaraju, Nasir Bhanpuri, David Miller, Ben Yuhas, Kecia Addison, Shihching Liu, Marilyn Powell and Rayid Ghani
LAK - Learning Analytics and Knowledge Conference, 2015.

[25] What's in a name? Understanding the Interplay between Titles, Content, and Communities in Social Media

Himabindu Lakkaraju, Julian McAuley, Jure Leskovec
ICWSM - International AAAI Conference on Weblogs and Social Media, 2013.
Featured in Time, Forbes, Phys.Org, Business Insider, New Scientist
Oral Presentation (Top 3%)

[24] Dynamic Multi-Relational Chinese Restaurant Process for Analyzing Influences on Users in Social Media
Himabindu Lakkaraju, Indrajit Bhattacharya, Chiranjib Bhattacharyya
ICDM - IEEE International Conference on Data Mining, 2012.
Oral Presentation (Top 8.6%)

[23] Attention prediction on social media brand pages
Himabindu Lakkaraju, Jitendra Ajmera
CIKM - ACM Conference on Information and Knowledge Management, 2011.

[22] Exploiting Coherence for the Simultaneous Discovery of Latent Facets and associated Sentiments
Himabindu Lakkaraju, Chiranjib Bhattacharyya, Indrajit Bhattacharya, Srujana Merugu
SDM - SIAM International Conference on Data Mining, 2011.
Best Paper Award

[21] TEM: A novel perspective to modeling content on microblogs
Himabindu Lakkaraju, Hyung-Il-Ahn
WWW - International World Wide Web Conference, 2011.

[20] Smart news feeds for social networks using scalable joint latent factor models
Himabindu Lakkaraju, Angshu Rai, Srujana Merugu
WWW - International World Wide Web Conference, 2011.

Selected Preprints, Working Papers, and Workshop Articles

[19] The Disagreement Problem in Explainable Machine Learning: A Practitioner's Perspective [PDF] (under review)
Satyapriya Krishna*, Tessa Han*, Alex Gu, Shahin Jabbari, Steven Wu, Himabindu Lakkaraju
Preliminary version presented at CHI Workshop on Trust and Reliance in Human-AI Teams, 2022; **Featured in Fortune Magazine**.

[18] In-Context Unlearning: Language Models as Few Shot Unlearners [PDF] (under review)
Martin Pawelczyk, Seth Neel, Himabindu Lakkaraju

[17] Are Large Language Models Post Hoc Explainers? [PDF] (under review)
Nicholas Kroeger, Dan Ley, Satyapriya Krishna, Chirag Agarwal, Himabindu Lakkaraju

[16] Certifying LLM Safety against Adversarial Prompting [PDF] (under review)
Aounon Kumar, Chirag Agarwal, Suraj Srinivas, Aaron Li, Soheil Feizi, Himabindu Lakkaraju

[15] Quantifying Uncertainty in Natural Language Explanations of Large Language Models [PDF] (under review)
Sree Harsha Tanneru, Chirag Agarwal, Himabindu Lakkaraju

[14] Accurate, Explainable, and Private Models: Providing Recourse While Minimizing Training Data Leakage [PDF] (under review)
Catherine Huang, Chelsea Swoopes, Christina Xiao, Jiaqi Ma, Himabindu Lakkaraju

[13] Efficient Estimation of the Local Robustness of Machine Learning Models [PDF] (under review)
Tessa Han, Suraj Srinivas, Himabindu Lakkaraju

[12] Analyzing chain-of-thought prompting in Large language models via gradient-based feature Attributions [PDF] (under review)
Skyler Wu, Eric Shen, Charumathi Badrinath, Jiaqi Ma, Himabindu Lakkaraju

[11] Rethinking Explainability as a Dialogue: A Practitioner's Perspective [PDF]
(under review)
Himabindu Lakkaraju, Dylan Slack, Yuxin Chen, Chenhao Tan, Sameer Singh
Preliminary version presented at NeurIPS Workshop on Human-Centered AI, 2022.

[10] On the Impact of Adversarially Robust Models on Algorithmic Recourse [PDF]
(under review)
Satyapriya Krishna, Chirag Agarwal, Himabindu Lakkaraju
Preliminary version presented at NeurIPS Workshop on Trustworthy and Socially Responsible ML, 2022.

[9] Rethinking Stability for Attribution-Based Explanations [PDF] (under review)
Chirag Agarwal, Nari Johnson, Martin Pawelczyk, Satyapriya Krishna, Eshika Saxena, Marinka Zitnik, Himabindu Lakkaraju
Preliminary version presented at ICLR PAIR2Struct Workshop, 2022.

[8] When Algorithms Explain Themselves: AI Adoption and Accuracy of Experts' Decisions (working paper)
Himabindu Lakkaraju, Chiara Farronato

[7] Can Model Explanations Help Reduce Biases in Real-World Decision Making?
(working paper)
Himabindu Lakkaraju, Sarah Tan

[6] Enforcing Right to Explanation: Bridging the Gaps between ML Research and Policy
(working paper)
Himabindu Lakkaraju, Jiaqi Ma

[5] On the Incompatibility between AI Regulatory Guidelines (working paper)
Paul Hamilton, Davor Ljubenkov, Jiaqi Ma, Himabindu Lakkaraju

[4] An Empirical Study of the Trade-offs between Interpretability and Fairness [PDF]
Shahin Jabbari, Han-Ching Ou, Himabindu Lakkaraju, Milind Tambe
ICML Workshop on Human Interpretability in Machine Learning, 2020

[3] Aspect Specific Sentiment Analysis using Hierarchical Deep Learning [PDF]
Himabindu Lakkaraju, Richard Socher, Christopher Manning
NIPS Workshop on Deep Learning and Representation Learning, 2014

Patents

[2] Extraction and Grouping of Feature Words
Chiranjib Bhattacharyya, Himabindu Lakkaraju, Sunil Aravindam, Kaushik Nath
US8484228 B2

[1] Enhancing knowledge bases using rich social media
Jitendra Ajmera, Shantanu Godbole, Himabindu Lakkaraju, Ashish Verma
US20130224714 A1

Advising & Mentoring	**Current Advisees**:	
	Suraj Srinivas, Postdoctoral Fellow, Harvard University	2022 - Present
	Aounon Kumar, Postdoctoral Fellow, Harvard University	2023 - Present
	Chirag Agarwal, Postdoctoral Fellow, Harvard University	2020 - Present
	Martin Pawelczyk, Postdoctoral Fellow, Harvard University	2023 - Present
	Tessa Han, PhD Student, Harvard University	2020 - Present
	Satyapriya Krishna, PhD Student, Harvard University	2021 - Present
	Dan Ley, PhD Student, Harvard University	2022 - Present
	Alex Oesterling, PhD Student, Harvard University	2022 - Present
	Usha Bhalla, PhD Student, Harvard University	2022 - Present
	Paul Hamilton, PhD Student, Harvard Univeristy	2023 - Present
	Elita Lobo, PhD Student, UMass Amherst	2023 - Present
	Sree Harsha Tanneru, Masters Student, Harvard University	2023 - Present
	Nikhil Nayak, Masters Student, Harvard University	2023 - Present
	Aaron Li, Masters Student, Harvard University	2023 - Present

Yanchen Liu, Masters Student, Harvard University   2023 - Present
Charu Badrinath, Undergrad, Harvard University   2023 - Present
Eric Shen, Undergrad, Harvard University   2023 - Present
Catherine Huang, Undergrad, Harvard University   2023 - Present
Christina Xiao, Undergrad, Harvard University   2023 - Present

Past Advisees and Interns:
Jiaqi Ma (Postdoc, Harvard University => Assistant Professor, UIUC)
Dylan Slack (PhD, UC Irvine => Research Scientist, Scale AI)
Sophie Hilgard (PhD, Harvard University => Research Scientist, Twitter)
Aditya Karan (MS, Harvard University => PhD Student, UIUC CS)
Kaivalya Rawal (MS, Harvard University => Research Fellow, Oxford University)
Alexis Ross (Undergraduate, Harvard University => PhD Student, MIT EECS)
Isha Puri (Undergraduate, Harvard University => PhD Student, MIT EECS)
Emily Jia (Undergraduate, Harvard University => Data Scientist, Figma)
Umang Bhatt (Research Intern, Harvard University => Assistant Professor, NYU CDS)
Harvineet Singh (Research Intern, Harvard University => Postdoc, UCSF/UC Berkeley)
Jessica Dai (Research Intern, Harvard University => PhD Student, UC Berkeley EECS)

Teaching Experience	Instructor, Explainable Artificial Intelligence Department of Computer Science, Harvard University (First ever full-fledged course on this topic)	2020 - 2023
	Instructor, Introduction to Data Science and Machine Learning Harvard Business School	2020 - 2023
	Instructor, A Short Course on Explainable Machine Learning Stanford Center for AI Safety	2022
	Instructor, Introduction to ML for Social Scientists Harvard Business School & Department of Computer Science	Spring 2020
	Instructor, Explainable and Accurate AI for High-Stakes Decision Making Harvard Online Analytics Program	2020 - 2022
	Guest Lecture, User Evaluations in Explainable Machine Learning UC Berkeley: Human-Centered AI Course	Spring 2023
	Guest Lecture, Explainable ML in the Era of Foundation Models Carnegie Mellon University: Trustworthy AI Course	Spring 2023
	Guest Lecture, Evaluating ML Models in the Presence of Unobservables Stanford University: Counterfactuals: The Science of What Ifs?	Spring 2021
	Guest Lecture, Explainable Machine Learning Harvard University: AI for Social Impact Course	Spring 2021
	Guest Lecture, Explainable Machine Learning Carnegie Mellon University: Advanced Introduction to Machine Learning Course	Autumn 2020
	Guest Lecture, Explainable Machine Learning in Practice Carnegie Mellon University: Human-AI Interaction Course	Autumn 2020
	Guest Lecture, Introduction to Data Science, Stanford Law School	Spring 2016
	Guest Lecture, Algorithms for Submodular Optimization Stanford University: Mining Massive Data Sets Course	Winter 2016
	Co-instructor, Introduction to Python Programming Stanford University: Girls Teaching Girls to Code (GTGTC) Initiative	Spring 2015
	Teaching Assistant for Stanford University: Mining Massive Data Sets Course Stanford University: Social & Information Network Analysis Course	Winter 2016 Autumn 2014

Indian Institute of Science: Machine Learning Course   Autumn 2010

Tutorials   Trustworthy Machine Learning in the Era of Foundation Models   ICML, FAccT, KDD 2023

Model Monitoring in Practice: Lessons Learned and Open Challenges   KDD, FAccT 2022

Explainable ML in the Wild: When Not to Trust Your Explanations   FAccT 2021

Explainable ML: Understanding the Limits and Pushing the Boundaries   CHIL 2021
(Invited Tutorial)

Explaining Machine Learning Predictions: NeurIPS 2020
State-of-the-art, Challenges, and Opportunities   AAAI 2021

Invited Talks   MIT Data Science Talk Series   2023
US Securities and Exchange Commission   2023
Google, Stanford, and UW Madison Workshop on Securing the Future of GenAI   2023
Yale and Google Joint Workshop on Theory and Practice of Foundation Models 2023
ICML Workshop on Interpretable ML in Healthcare   2023
ICML Workshop on Counterfactuals in Minds and Machines   2023
ICLR Workshop on Trustworthy & Reliable Large-Scale Machine Learning Models 2023
RSS Workshop on Safe Autonomy   2023
Mind and Machine Intelligence Summit, UC Santa Barbara   2023
Cornell University and Weill Cornell Medicine   2023
Kavli Frontiers of Science Symposium   2023
Cohere AI   2023
Keynote at AAAI Workshop on Representation Learning for   2023
Responsible Human-Centric AI
Keynote at AAAI Workshop on Deployable AI   2023
INFORMS Annual Meeting   2016 - 2023
NeurIPS Workshop on Women in Machine Learning (WiML)   2022
NeurIPS Workshop on Machine Learning for Health (ML4H)   2022
ICLR Workshop on Privacy, Accountability, Interpretability, Robustness,   2022
Reasoning on Structured Data
CVPR Workshop on Explainable AI for Computer Vision   2022
Keynote at WWW Workshop on Explainable AI in Health   2022
ECCV Workshop on Adversarial Robustness in the Real World   2022
Panel Discussion on AI and the Economy, Jointly Organized by   2022
U.S. Department of Commerce, NIST, Stanford HAI, and the FinRegLab
Simons Institute (Berkeley) Workshop on Societal Considerations and Applications 2022
Stanford Center for AI Safety Workshop on Explainable AI   2022
Stanford Human-Centered Artificial Intelligence (HAI) Conference   2022
Stanford Digital Econ Seminar   2022
MIT Initiative on the Digital Economy (IDE) Seminar Series   2022
Harvard Data Science Initiative's Annual Conference   2022
Berkman Klein Center, Harvard University   2022
Amazon Alexa Rising Star Speaker Series   2022
University of Southern California   2022
Fireside Chat on Explainability, Fiddler AI   2022
Keynote at ACM CIKM Conference   2021
NIST AI Risk Management Framework Workshop   2021
Pinterest **Distinguished Lecture**   2021
NeurIPS Workshop on Algorithmic Fairness through the Lens of   2021
Causality and Robustness
NeurIPS Workshop on Explainable AI Approaches for Debugging and Diagnosis 2021
NeurIPS Workshop on Human and Machine Decisions   2021
Keynote at ICML Workshop on Interpretable ML in Healthcare   2021
Keynote at KDD Workshop on ML in finance   2021
AI for Good Summit organized by International Telecommunications Union &   2021

the United Nations

Keynote at CVPR Workshop on Responsible Computer Vision	2021
Keynote at ICLR Workshop on Responsible AI	2021
Keynote at ASPLOS Workshop on Systems Architecture for Robust, Safe, and Resilient Software	2021
Keynote at MLSys Workshop on Personalized Recommender Systems & Algorithms	2021
University of Cambridge	2021
Neurosym Webinar Series, Jointly Organized by UPenn, MIT, Caltech, and Stanford	2021
Voices of Data Science, UMass Amherst	2021
Max Planck Symposium on Computing and Society	2021
Keynote at CVPR Workshop on Fair, Data-Efficient and Trusted Computer Vision	2020
Keynote at MICCAI Workshop on Interpretability in Medical Imaging	2020
ETH - Center for Law and Economics, Zurich	2020
University of Michigan, Ann Arbor	2019
Harvard CRCS Seminar, Cambridge	2019
AI World Conference & Expo, Cambridge	2019
EmTech MIT Conference, Cambridge	2019
Google DeepMind Annual Summit, Cambridge	2019
Women in Machine Learning Workshop, Boston	2019
ICLR Workshop on Safe Machine Learning, New Orleans	2019
Harvard Data Science Conference, Cambridge	2018
South Park Commons, San Francisco	2018
Computer Science Departmental Seminars at Carnegie Mellon University, UIUC, Harvard University, Georgia Tech, Yale University, UC San Diego, USC, UCLA, UC Irvine, Duke University, Brown University, University of Michigan, University of Maryland	2018
Machine Learning Departmental Seminar at Carnegie Mellon University	2018
Operations Research Departmental Seminars at Columbia University, Cornell University, Princeton University	2018
NYU Stern School of Business, New York	2018
MIT Sloan School of Management, Cambridge	2018
Harvard Business School, Boston	2018
UC Berkeley School of Public Health, San Francisco	2018
Microsoft Research, Redmond	2017, 2018
IBM Thomas J. Watson Research Center, New York	2017
Machine Learning Seminar at Duke University, Durham	2017
Keynote at ICML Workshop on Automatic Machine Learning, Sydney, Australia	2017
Stanford Biomedical Data Science Lecture Series, Palo Alto	2017
Stanford Symbolic Systems Coffee Chat Series, Palo Alto	2017
Stanford Data Science Workshop, Palo Alto	2017
Rising Stars Workshop in EECS, Pittsburgh	2016
CodeX Center, Stanford Law School, Palo Alto	2016
KDD Workshop on Data Science for Social Good, New York	2014
University of Chicago Computation Institute, Chicago	2014
Grace Hopper India Chapter, Bangalore, India	2011

Community Service

Co-Founder & Chair: Trustworthy ML Initiative	2020 - Present

We launched this initiative to enable easy access to resources on trustworthy ML, to showcase and promote the work of researchers from underrepresented groups, and to build a community of researchers and practitioners working on the topic.

Panelist and Reviewer:	2020 - Present

4 National Science Foundation (NSF) Review Panels,
Directorate for Computer and Information Science and Engineering (CISE)

Co-Chair:

NeurIPS Workshop on Regulatable Machine Learning	2023
NeurIPS Workshop on Explainable Artificial Intelligence	2023
KDD Trustworthy AI Day	2022

ICML Workshop on New Frontiers in Adversarial Machine Learning	2022
KDD Deep Learning Day	2021
ICML Workshop on Algorithmic Recourse	2021
ELLIS Human-Centric Machine Learning Workshop	2021
Session on Trustworthy Machine Learning at INFORMS	2020
Session on Fairness in Machine Learning at INFORMS	2019
ICLR Workshop on Debugging Machine Learning Models	2019
Workshop for spreading awareness about STEM fields among middle school girls	2016
Stanford's Girls Teaching Girls To Code (GTGTC)	2015
Grace Hopper India Conference	2011

Sponsorship Chair:

FAccT - ACM Conference on Fairness, Accountability, and Transparency	2023

Tutorial Chair:

WSDM - ACM Conference on Web Search and Data Mining	2024

Area Chair:

NeurIPS - *Advances in Neural Information Processing Systems*	2019 - 2023
ICLR - *International Conference on Learning Representations*	2020 - 2023
AISTATS - *International Conference on Artificial Intelligence and Statistics*	2021 - 2023
ICML - *International Conference on Machine Learning*	2019 - 2023

Program Committee:

AISTATS - *International Conference on Artificial Intelligence and Statistics*	2019 - 2020
FAccT - *ACM Conference on Fairness, Accountability, and Transparency*	2019 - 2020
AAAI - *AAAI International Conference on Artificial Intelligence*	2019
ICML - *International Conference on Machine Learning*	2018
ICLR - *International Conference on Learning Representations*	2018 - 2019
IJCAI - *International Joint Conference on Artificial Intelligence*	2018 - 2019
WWW - *International World Wide Web Conference*	2017 - 2018
NIPS - *Advances in Neural Information Processing Systems*	2016 - 2017
KDD - *ACM SIGKDD Conference on Knowledge Discovery and Data Mining*	2015 - 2017
CIKM - *ACM Conference on Information and Knowledge Management*	2011, 2017
SDM - *SIAM International Conference on Data Mining*	2015
UAI - *Conference on Uncertainty in Artificial Intelligence*	2011
AAAI - *AAAI conference on Artificial Intelligence*	2011

Journal Reviewing and Editing:

Frontiers in Big Data (Associate Editor)	2021 - 2023
JMLR - *Journal of Machine Learning Research*	2020 - 2023
MS - *Management Science*	2021 - 2023
OR - *Operations Research*	2021 - 2023
TWEB - *ACM Transactions on the Web*	2017
PLOS ONE - *Public Library of Science ONE*	2017
TKDD - *ACM Transactions on Knowledge Discovery from Data*	2016
TKDE - *IEEE Transactions on Knowledge and Data Engineering*	2015

Other:

Member, Faculty Hiring Committee, Harvard University	2020 - 2023
Member, Ph.D. Student Selection Committee, Harvard University	2020 - 2023
Member, Ph.D. Student Selection Committee, Stanford University	2016

Selected Media Coverage

TIME: Chuck Schumer wants AI to be explainable. It's harder than it sounds
Fortune: Explainable AI & The Disagreement Problem
Harvard Business Review: The AI transparency paradox
MIT Technology Review: How to upgrade judges with machine learning
Harvard Business Review: Solving social problems with machine learning
The New York Times: Even Imperfect Algorithms Can Improve the Criminal Justice System
VentureBeat: Confidence, uncertainty, and trust in AI affect how humans make decisions
Wired: This Agency Wants to Figure Out Exactly How Much You Trust AI

Bloomberg Technology: Researchers combat gender and racial bias in AI
Forbes: How to craft the perfect Reddit posting
Time: How to succeed on Reddit
Business Insider: How to execute the perfect Reddit submission
Phys.org: Stanford Trio explore success formula for Reddit posts
International Business Times: The secret to what makes something go viral
New Scientist: Things that make a meme explode
The Verge: The math behind successful Reddit submissions
ACM TechNews: Stanford trio explore success formula for Reddit posts